Rule #4

Drain the Shallows

In the summer of 2007, the software company 37signals (now called Basecamp) launched an experiment: They shortened their workweek from five days to four. Their employees seemed to accomplish the same amount of work with one less day, so they made this change permanent: Every year, from

May through October, 37signals employees work only Monday to Thursday (with the exception of customer support, which still operates the full week). As company cofounder Jason Fried quipped in a blog post about the decision: "People should enjoy the weather in the summer."

It didn't take long before the grumbles began in the business press. A few months after Fried announced his company's decision to make four-day weeks permanent, journalist Tara Weiss wrote a critical piece for *Forbes* titled "Why a Four-Day Work Week Doesn't Work." She summarized her problem with this strategy as follows:

Packing 40 hours into four days isn't necessarily an efficient way to work. Many people find that eight hours are tough enough; requiring them to stay for an extra two could cause morale and productivity to decrease.

Fried was quick to respond. In a blog post titled "Forbes Misses the Point of the 4-Day Work Week," he begins by agreeing with Weiss's premise that it *would* be stressful for employees to cram forty hours of effort into four days. But, as he clarifies, that's not what he's suggesting. "The point of the 4-day work week is about *doing less work*," he

writes. "It's not about four 10-hour days… it's about four normalish 8-hour days."

This might seem confusing at first. Fried earlier claimed that his employees get just as much done in four days as in five days. Now, however, he's claiming that his employees are working fewer hours. How can both be true? The difference, it turns out, concerns the role of shallow work. As Fried expands:

> *Very few people work even 8 hours a day. You're lucky if you get a few good hours in between all the meetings, interruptions, web surfing, office politics, and*

personal business that permeate the typical workday.

Fewer official working hours helps squeeze the fat out of the typical workweek. Once everyone has less time to get their stuff done, they respect that time even more. People become stingy with their time and that's a good thing. They don't waste it on things that just don't matter. When you have fewer hours you usually spend them more wisely.

In other words, the reduction in the 37signals workweek disproportionately eliminated shallow as compared to deep

work, and because the latter was left largely untouched, the important stuff continued to get done. The shallow stuff that can seem so urgent in the moment turned out to be unexpectedly dispensable.

A natural reaction to this experiment is to wonder what would happen if 37signals had gone one step further. If eliminating hours of shallow work had little impact on the results produced, what would happen if they not only eliminated shallow work, but then replaced this newly recovered time with more deep work? Fortunately for our curiosity, the company soon put this bolder idea to the test as well.

Fried had always been interested in the policies of technology companies like Google that gave their employees 20 percent of their time to work on self-directed projects. While he liked this idea, he felt that carving one day out of an otherwise busy week was not enough to support the type of unbroken deep work that generates true breakthroughs. "I'd take 5 days in a row over 5 days spread out over 5 weeks," he explained. "So our theory is that we'll see better results when people have a long stretch of uninterrupted time."

To test this theory, 37signals implemented something radical: The company gave its employees the *entire*

month of June off to work deeply on their own projects. This month would be a period free of any shallow work obligations—no status meetings, no memos, and, blessedly, no PowerPoint. At the end of the month, the company held a "pitch day" in which employees pitched the ideas they'd been working on. Summarizing the experiment in an *Inc.* magazine article, Fried dubbed it a success. The pitch day produced two projects that were soon put into production: a better suite of tools for handling customer support and a data visualization system that helps the company understand how their customers use their products. These

projects are predicted to bring substantial value to the company, but they almost certainly would *not* have been produced in the absence of the unobstructed deep work time provided to the employees. To tease out their potential required dozens of hours of unimpeded effort.

"How can we afford to put our business on hold for a month to 'mess around' with new ideas?" Fried asked rhetorically. "How can we afford not to?"

37signals' experiments highlight an important reality: The shallow work that increasingly dominates the time and

attention of knowledge workers is less vital than it often seems in the moment. For most businesses, if you eliminated significant amounts of this shallowness, their bottom line would likely remain unaffected. And as Jason Fried discovered, if you not only eliminate shallow work, but also replace this recovered time with more of the deep alternative, not only will the business continue to function; it can become *more* successful.

This rule asks you to apply these insights to your personal work life. The strategies that follow are designed to help you ruthlessly identify the shallowness in your current schedule,

then cull it down to minimum levels—leaving more time for the deep efforts that ultimately matter most.

Before diving into the details of these strategies, however, we should first confront the reality that there's a limit to this anti-shallow thinking. The value of deep work vastly outweighs the value of shallow, but this doesn't mean that you must quixotically pursue a schedule in which *all* of your time is invested in depth. For one thing, a nontrivial amount of shallow work is needed to maintain most knowledge work jobs. You might be able to avoid checking your e-mail every ten minutes, but you won't likely last long if you *never* respond to

important messages. In this sense, we should see the goal of this rule as taming shallow work's footprint in your schedule, not eliminating it.

Then there's the issue of cognitive capacity. Deep work is exhausting because it pushes you toward the limit of your abilities. Performance psychologists have extensively studied how much such efforts can be sustained by an individual in a given day.* In their seminal paper on deliberate practice, Anders Ericsson and his collaborators survey these studies. They note that for someone new to such practice (citing, in particular, a child in the early stages of developing an expert-level skill), an

hour a day is a reasonable limit. For those familiar with the rigors of such activities, the limit expands to something like four hours, but rarely more.

The implication is that once you've hit your deep work limit in a given day, you'll experience diminishing rewards if you try to cram in more. Shallow work, therefore, doesn't become dangerous until after you add enough to begin to crowd out your bounded deep efforts for the day. At first, this caveat might seem optimistic. The typical workday is eight hours. The most adept deep thinker cannot spend more than four of these hours in a state of true depth. It follows that you can safely spend half the day

wallowing in the shallows without adverse effect. The danger missed by this analysis is how *easily* this amount of time can be consumed, especially once you consider the impact of meetings, appointments, calls, and other scheduled events. For many jobs, these time drains can leave you with surprisingly little time left for solo work.

My job as a professor, for example, is traditionally less plagued by such commitments, but even so, they often take large chunks out of my time, especially during the academic year. Turning to a random day in my calendar from the previous semester (I'm writing this during a quiet summer month), for

example, I see I had a meeting from eleven to twelve, another from one to two thirty, and a class to teach from three to five. My eight-hour workday in this example is already reduced by four hours. Even if I squeezed all remaining shallow work (e-mails, tasks) into a single half hour, I'd still fall short of the goal of four hours of daily deep work. Put another way, even though we're not capable of spending a full day in a state of blissful depth, this reality shouldn't reduce the urgency of reducing shallow work, as the typical knowledge workday is more easily fragmented than many suspect.

To summarize, I'm asking you to treat

shallow work with suspicion because its damage is often vastly underestimated and its importance vastly overestimated. This type of work is inevitable, but you must keep it confined to a point where it doesn't impede your ability to take full advantage of the deeper efforts that ultimately determine your impact. The strategies that follow will help you act on this reality.

Schedule Every Minute of Your Day

If you're between the ages of twenty-five and thirty-four years old and live in

Britain, you likely watch more television than you realize. In 2013, the British TV licensing authority surveyed television watchers about their habits. The twenty-five-to thirty-four-year-olds taking the survey estimated that they spend somewhere between fifteen and sixteen hours per week watching TV. This sounds like a lot, but it's actually a significant underestimate. We know this because when it comes to television-watching habits, we have access to the ground truth. The Broadcasters' Audience Research Board (the British equivalent of the American Nielsen Company) places meters in a representative sample of households.

These meters record, without bias or wishful thinking, exactly how much people *actually* watch. The twenty-five-to thirty-four-year-olds who thought they watched fifteen hours a week, it turns out, watch more like twenty-eight hours.

This bad estimate of time usage is not unique to British television watching. When you consider different groups self-estimating different behaviors, similar gaps stubbornly remain. In a *Wall Street Journal* article on the topic, business writer Laura Vanderkam pointed out several more such examples. A survey by the National Sleep Foundation revealed that Americans think they're sleeping, on average, somewhere around

seven hours a night. The American Time Use Survey, which has people actually measure their sleep, corrected this number to 8.6 hours. Another study found that people who claimed to work sixty to sixty-four hours per week were actually averaging more like forty-four hours per week, while those claiming to work more than seventy-five hours were actually working less than fifty-five.

These examples underscore an important point: We spend much of our day on autopilot—not giving much thought to what we're doing with our time. *This is a problem.* It's difficult to prevent the trivial from creeping into every corner of your schedule if you

don't face, without flinching, your current balance between deep and shallow work, and then adopt the habit of pausing before action and asking, "What makes the most sense right now?" The strategy described in the following paragraphs is designed to force you into these behaviors. It's an idea that might seem extreme at first but will soon prove indispensable in your quest to take full advantage of the value of deep work: *Schedule every minute of your day*.

Here's my suggestion: At the beginning of each workday, turn to a new page of lined paper in a notebook you dedicate to this purpose. Down the left-hand side

of the page, mark every other line with an hour of the day, covering the full set of hours you typically work. Now comes the important part: Divide the hours of your workday into *blocks* and assign activities to the blocks. For example, you might block off nine a.m. to eleven a.m. for writing a client's press release. To do so, actually draw a box that covers the lines corresponding to these hours, then write "press release" inside the box. Not every block need be dedicated to a work task. There might be time blocks for lunch or relaxation breaks. To keep things reasonably clean, the minimum length of a block should be thirty minutes (i.e., one line on your

page). This means, for example, that instead of having a unique small box for each small task on your plate for the day —*respond to boss's e-mail, submit reimbursement form, ask Carl about report*—you can batch similar things into more generic *task blocks*. You might find it useful, in this case, to draw a line from a task block to the open right-hand side of the page where you can list out the full set of small tasks you plan to accomplish in that block.

When you're done scheduling your day, every minute should be part of a block. You have, in effect, given every minute of your workday a job. Now as you go through your day, use this

schedule to guide you.

It's here, of course, that most people will begin to run into trouble. Two things can (and likely will) go wrong with your schedule once the day progresses. The first is that your estimates will prove wrong. You might put aside two hours for writing a press release, for example, and in reality it takes two and a half hours. The second problem is that you'll be interrupted and new obligations will unexpectedly appear on your plate. These events will also break your schedule.

This is okay. If your schedule is disrupted, you should, at the next available moment, take a few minutes to

create a revised schedule for the time that remains in the day. You can turn to a new page. You can erase and redraw blocks. Or do as I do: Cross out the blocks for the remainder of the day and create new blocks to the right of the old ones on the page (I draw my blocks skinny so I have room for several revisions). On some days, you might rewrite your schedule half a dozen times. Don't despair if this happens. Your goal is not to stick to a given schedule at all costs; it's instead to maintain, at all times, a thoughtful say in what you're doing with your time going forward—even if these decisions are reworked again and again as the day

unfolds.

If you find that schedule revisions become overwhelming in their frequency, there are a few tactics that can inject some more stability. First, you should recognize that *almost definitely* you're going to underestimate at first how much time you require for most things. When people are new to this habit, they tend to use their schedule as an incarnation of wishful thinking—a best-case scenario for their day. Over time, you should make an effort to accurately (if not somewhat conservatively) predict the time tasks will require.

The second tactic that helps is the use

of *overflow conditional* blocks. If you're not sure how long a given activity might take, block off the expected time, then follow this with an additional block that has a split purpose. If you need more time for the preceding activity, use this additional block to keep working on it. If you finish the activity on time, however, have an alternate use already assigned for the extra block (for example, some nonurgent tasks). This allows unpredictability in your day without requiring you to keep changing your schedule on paper. For example, returning to our press release example, you might schedule two hours for writing the press release, but then follow it by

an additional hour block that you can use to keep writing the release, if needed, but otherwise assign to catching up with e-mail.

The third tactic I suggest is to be liberal with your use of task blocks. Deploy many throughout your day and make them longer than required to handle the tasks you plan in the morning. Lots of things come up during the typical knowledge worker's day: Having regularly occurring blocks of time to address these surprises keeps things running smoothly.

Before leaving you to put this strategy in practice, I should address a common

objection. In my experience pitching the values of daily schedules, I've found that many people worry that this level of planning will become burdensomely restrictive. Here, for example, is part of a comment from a reader named Joseph on a blog post I wrote on this topic:

> *I think you far understate the role of uncertainty... I [worry about] readers applying these observations too seriously, to the point of an obsessive (and unhealthy) relationship with one's schedule that seems to exaggerate the importance of minute-counting over getting-*

lost-in-activities, which if we're talking about artists is often the only really sensible course of action.

I understand these concerns, and Joseph is certainly not the first to raise them. Fortunately, however, they're also easily addressed. In my own daily scheduling discipline, in addition to regularly scheduling significant blocks of time for speculative thinking and discussion, I maintain a rule that if I stumble onto an important insight, then this is a perfectly valid reason to ignore the rest of my schedule for the day (with the exception, of course, of things that

cannot be skipped). I can then stick with this unexpected insight until it loses steam. At this point, I'll step back and rebuild my schedule for any time that remains in the day.

In other words, I not only allow spontaneity in my schedule; I encourage it. Joseph's critique is driven by the mistaken idea that the goal of a schedule is to force your behavior into a rigid plan. This type of scheduling, however, isn't about constraint—it's instead about thoughtfulness. It's a simple habit that forces you to continually take a moment throughout your day and ask: "What makes sense for me to do with the time that remains?" It's the habit of asking

that returns results, not your unyielding fidelity to the answer.

I would go so far as to argue that someone following this combination of comprehensive scheduling and a willingness to adapt or modify the plan as needed will likely experience *more* creative insights than someone who adopts a more traditionally "spontaneous" approach where the day is left open and unstructured. Without structure, it's easy to allow your time to devolve into the shallow—e-mail, social media, Web surfing. This type of shallow behavior, though satisfying in the moment, is not conducive to creativity. With structure, on the other

hand, you can ensure that you regularly schedule blocks to grapple with a new idea, or work deeply on something challenging, or brainstorm for a fixed period—the type of commitment more likely to instigate innovation. (Recall, for example, the discussion in Rule #1 about the rigid rituals followed by many great creative thinkers.) And because you're willing to abandon your plan when an innovative idea arises, you're just as well suited as the distracted creative to follow up when the muse strikes.

To summarize, the motivation for this strategy is the recognition that a deep

work habit requires you to treat your time with respect. A good first step toward this respectful handling is the advice outlined here: Decide in advance what you're going to do with every minute of your workday. It's natural, at first, to resist this idea, as it's undoubtedly easier to continue to allow the twin forces of internal whim and external requests to drive your schedule. But you must overcome this distrust of structure if you want to approach your true potential as someone who creates things that matter.

Quantify the Depth of Every

Activity

An advantage of scheduling your day is that you can determine how much time you're actually spending in shallow activities. Extracting this insight from your schedules, however, can become tricky in practice, as it's not always clear exactly how shallow you should consider a given task. To expand on this challenge, let's start by reminding ourselves of the formal definition of shallow work that I introduced in the introduction:

Shallow Work: Noncognitively demanding, logistical-style tasks, often

performed while distracted. These efforts tend not to create much new value in the world and are easy to replicate.

Some activities clearly satisfy this definition. Checking e-mail, for example, or scheduling a conference call, is unquestionably shallow in nature. But the classification of other activities can be more ambiguous. Consider the following tasks:

- **Example #1**: Editing a draft of an academic article that you and a collaborator will soon submit to a journal.
- **Example #2**: Building a

PowerPoint presentation about this quarter's sales figures.

- **Example #3**: Attending a meeting to discuss the current status of an important project and to agree on the next steps.

It's not obvious at first how to categorize these examples. The first two describe tasks that can be quite demanding, and the final example seems important to advance a key work objective. The purpose of this strategy is to give you an accurate metric for resolving such ambiguity—providing you with a way to make clear and consistent decisions about where given

work tasks fall on the shallow-to-deep scale. To do so, it asks that you evaluate activities by asking a simple (but surprisingly illuminating) question:

> How long would it take (in months) to train a smart recent college graduate with no specialized training in my field to complete this task?

To illustrate this approach, let's apply this question to our examples of ambiguous tasks.

- **Analyzing Example #1**: To properly edit an academic paper

requires that you understand the nuances of the work (so you can make sure it's being described precisely) and the nuances of the broader literature (so you can make sure it's being cited properly). These requirements require cutting-edge knowledge of an academic field—a task that in the age of specialization takes years of diligent study at the graduate level and beyond. When it comes to this example, the answer to our question would therefore be quite large, perhaps on the scale of fifty to seventy-five months.

- **Analyzing Example #2**: The

second example doesn't fare so well by this analysis. To create a PowerPoint presentation that describes your quarterly sales requires three things: first, knowledge of how to make a PowerPoint presentation; second, an understanding of the standard format of these quarterly performance presentations within your organization; and third, an understanding of what sales metrics your organization tracks and how to convert them into the right graphs. The hypothetical college graduate imagined by our question, we can assume, would already know how to use PowerPoint, and learning the

standard format for your organization's presentations shouldn't require more than a week. The real question, therefore, is how long it takes a bright college graduate to understand the metrics you track, where to find the results, and how to clean those up and translate them into graphs and charts that are appropriate for a slide presentation. This isn't a trivial task, but for a bright college grad it wouldn't require more than an additional month or so of training —so we can use two months as our conservative answer.

- **Analyzing Example #3**: Meetings

can be tricky to analyze. They can seem tedious at times but they're often also presented as playing a key role in your organization's most important activities. The method presented here helps cut through this veneer. How long would it take to train a bright recent college graduate to take your place in a planning meeting? He or she would have to understand the project well enough to know its milestones and the skills of its participants. Our hypothetical grad might also need some insight into the interpersonal dynamics and the reality of how such projects are executed at the organization. At this point, you might

wonder if this college grad would also need a deep expertise in the topic tackled by the project. For a planning meeting—probably not. Such meetings rarely dive into substantive content and tend to feature a lot of small talk and posturing in which participants try to make it seem like they're committing to a lot without actually having to commit. Give a bright recent graduate three months to learn the ropes and he or she could take your place without issue in such a gabfest. So we'll use three months as our answer.

This question is meant as a thought

experiment (I'm not going to ask you to actually *hire* a recent college graduate to take over tasks that score low). But the answers it provides will help you objectively quantify the shallowness or depth of various activities. If our hypothetical college graduate requires many months of training to replicate a task, then this indicates that the task leverages hard-won expertise. As argued earlier, tasks that leverage your expertise tend to be deep tasks and they can therefore provide a double benefit: They return more value per time spent, and they stretch your abilities, leading to improvement. On the other hand, a task that our hypothetical college graduate

can pick up quickly is one that does not leverage expertise, and therefore it can be understood as shallow.

What should you do with this strategy? Once you know where your activities fall on the deep-to-shallow scale, bias your time toward the former. When we reconsider our case studies, for example, we see that the first task is something that you would want to prioritize as a good use of time, while the second and third are activities of a type that should be minimized—they might feel productive, but their return on (time) investment is measly.

Of course, how one biases away from shallow and toward depth is not

always obvious—even after you know how to accurately label your commitments. This brings us to the strategies that follow, which will provide specific guidance on how to accomplish this tricky goal.

Ask Your Boss for a Shallow Work Budget

Here's an important question that's rarely asked: *What percentage of my time should be spent on shallow work?* This strategy suggests that you ask it. If you have a boss, in other words, have a conversation about this question. (You'll

probably have to first define for him or her what "shallow" and "deep" work means.) If you work for yourself, ask *yourself* this question. In both cases, settle on a specific answer. Then—and this is the important part—try to stick to this budget. (The strategies that precede and follow this one will help you achieve this goal.)

For most people in most non-entry-level knowledge work jobs, the answer to the question will be somewhere in the 30 to 50 percent range (there's a psychological distaste surrounding the idea of spending the *majority* of your time on unskilled tasks, so 50 percent is a natural upper limit, while at the same

time most bosses will begin to worry that if this percentage gets too much lower than 30 percent you'll be reduced to a knowledge work hermit who thinks big thoughts but never responds to e-mails).

Obeying this budget will likely require changes to your behavior. You'll almost certainly end up forced into saying no to projects that seem infused with shallowness while also more aggressively reducing the amount of shallowness in your existing projects. This budget might lead you to drop the need for a weekly status meeting in preference for results-driven reporting ("let me know when you've made

significant progress; then we'll talk"). It might also lead you to start spending more mornings in communication isolation or decide it's not as important as you once thought to respond quickly and in detail to every cc'd e-mail that crosses your inbox.

These changes are all positive for your quest to make deep work central to your working life. On the one hand, they don't ask you to abandon your core shallow obligations—a move that would cause problems and resentment—as you're still spending a lot of time on such efforts. On the other hand, they do force you to place a hard limit on the amount of less urgent obligations you

allow to slip insidiously into your schedule. This limit frees up space for significant amounts of deep effort on a consistent basis.

The reason why these decisions should start with a conversation with your boss is that this agreement establishes implicit support from your workplace. If you work for someone else, this strategy provides cover when you turn down an obligation or restructure a project to minimize shallowness. You can justify the move because it's necessary for you to hit your prescribed target mix of work types. As I discussed in Chapter 2, part of the reason shallow work persists in large

quantities in knowledge work is that we rarely see the total impact of such efforts on our schedules. We instead tend to evaluate these behaviors one by one in the moment—a perspective from which each task can seem quite reasonable and convenient. The tools from earlier in this rule, however, allow you to make this impact explicit. You can now confidently say to your boss, "This is the exact percentage of my time spent last week on shallow work," and force him or her to give explicit approval for that ratio. Faced with these numbers, and the economic reality they clarify (it's incredibly wasteful, for example, to pay a highly trained professional to send e-

mail messages and attend meetings for thirty hours a week), a boss will be led to the natural conclusion that you *need* to say no to some things and to streamline others—even if this makes life less convenient for the boss, or for you, or for your coworkers. Because, of course, in the end, a business's goal is to generate value, not to make sure its employees' lives are as easy as possible.

If you work for yourself, this exercise will force you to confront the reality of how little time in your "busy" schedule you're actually producing value. These hard numbers will provide you the confidence needed to start

scaling back on the shallow activities that are sapping your time. Without these numbers, it's difficult for an entrepreneur to say no to any opportunity that *might* generate some positive return. "I have to be on Twitter!," "I have to maintain an active Facebook presence!," "I have to tweak the widgets on my blog!," you tell yourself, because when considered in isolation, to say no to any one of these activities seems like you're being lazy. By instead picking and sticking with a shallow-to-deep ratio, you can replace this guilt-driven unconditional acceptance with the more healthy habit of trying to get the most out of the time you put aside for shallow

work (therefore still exposing yourself to many opportunities), but keeping these efforts constrained to a small enough fraction of your time and attention to enable the deep work that ultimately drives your business forward.

Of course, there's always the possibility that when you ask this question the answer is stark. No boss will explicitly answer, "One hundred percent of your time should be shallow!" (unless you're entry level, at which point you need to delay this exercise until you've built enough skills to add deep efforts to your official work responsibilities), but a boss might reply, in so many words, "as much shallow

work as is needed for you to promptly do whatever we need from you at the moment." In this case, the answer is still useful, as it tells you that this isn't a job that supports deep work, and a job that doesn't support deep work is not a job that can help you succeed in our current information economy. You should, in this case, thank the boss for the feedback, and then promptly start planning how you can transition into a new position that values depth.

Finish Your Work by Five Thirty

In the seven days preceding my first writing these words, I participated in sixty-five different e-mail conversations. Among these sixty-five conversations, I sent exactly five e-mails after five thirty p.m. The immediate story told by these statistics is that, with few exceptions, I don't send e-mails after five thirty. But given how intertwined e-mail has become with work in general, there's a more surprising reality hinted by this behavior: I don't *work* after five thirty p.m.

I call this commitment *fixed-schedule productivity*, as I fix the firm goal of not working past a certain time, then work backward to find productivity

strategies that allow me to satisfy this declaration. I've practiced fixed-schedule productivity happily for more than half a decade now, and it's been crucial to my efforts to build a productive professional life centered on deep work. In the pages ahead, I will try to convince you to adopt this strategy as well.

Let me start my pitch for fixed-schedule productivity by first noting that, according to conventional wisdom, in the academic world I inhabit this tactic should fail. Professors—especially junior professors—are notorious for adopting grueling schedules that extend

into the night and through weekends. Consider, for example, a blog post published by a young computer science professor whom I'll call "Tom." In this post, which Tom wrote in the winter of 2014, he replicates his schedule for a recent day in which he spent twelve hours at his office. This schedule includes five different meetings and three hours of "administrative" tasks, which he describes as "tending to bushels of e-mails, filling out bureaucratic forms, organizing meeting notes, planning future meetings." By his estimation, he spent only *one and a half* out of the twelve total hours sitting in his office tackling "real" work, which he

defines as efforts that make progress toward a "research deliverable." It's no wonder that Tom feels coerced into working well beyond the standard workday. "I've already accepted the reality that I'll be working on weekends," he concludes in another post. "Very few junior faculty can avoid such a fate."

And yet, *I have*. Even though I don't work at night and rarely work on weekends, between arriving at Georgetown in the fall of 2011 and beginning work on this chapter in the fall of 2014, I've published somewhere around twenty peer-reviewed articles. I also won two competitive grants,

published one (nonacademic) book, and have almost finished writing another (which you're reading at the moment). All while avoiding the grueling schedules deemed necessary by the Toms of the world.

What explains this paradox? We can find a compelling answer in a widely disseminated article published in 2013 by an academic further along in her career, and far more accomplished than I: Radhika Nagpal, the Fred Kavli Professor of Computer Science at Harvard University. Nagpal opens the article by claiming that much of the stress suffered by tenure-track professors is self-imposed. "Scary

myths and scary data abound about life as a tenure-track faculty at an 'R1' [research-focused] university," she begins, before continuing to explain how she finally decided to disregard the conventional wisdom and instead "deliberately… do specific things to preserve my happiness." This deliberate effort led Nagpal to enjoy her pre-tenure time "tremendously."

Nagpal goes on to detail several examples of these efforts, but there's one tactic in particular that should sound familiar. As Nagpal admits, early in her academic career she found herself trying to cram work into every free hour between seven a.m. and midnight

(because she has kids, this time, especially in the evening, was often severely fractured). It didn't take long before she decided this strategy was unsustainable, so she set a limit of fifty hours a week and worked backward to determine what rules and habits were needed to satisfy this constraint. Nagpal, in other words, deployed fixed-schedule productivity.

We know this strategy didn't hurt her academic career, as she ended up earning tenure on schedule and then jumping to the full professor level after only three additional years (an impressive ascent). How did she pull this off? According to her article, one of

the main techniques for respecting her hour limit was to set drastic quotas on the major sources of *shallow* endeavors in her academic life. For example, she decided she would travel only five times per year for any purpose, as trips can generate a surprisingly large load of urgent shallow obligations (from making lodging arrangements to writing talks). Five trips a year may still sound like a lot, but for an academic it's light. To emphasize this point, note that Matt Welsh, a former colleague of Nagpal in the Harvard computer science department (he now works for Google) once wrote a blog post in which he claimed it was typical for junior faculty

to travel twelve to twenty-four times a year. (Imagine the shallow efforts Nagpal avoided in sidestepping an extra ten to fifteen trips!) The travel quota is just one of several tactics that Nagpal used to control her workday (she also, for example, placed limits on the number of papers she would review per year), but what all her tactics shared was a commitment to ruthlessly capping the shallow while protecting the deep efforts—that is, original research—that ultimately determined her professional fate.

Returning to my own example, it's a similar commitment that enables me to succeed with fixed scheduling. I, too, am

incredibly cautious about my use of the most dangerous word in one's productivity vocabulary: "yes." It takes a lot to convince me to agree to something that yields shallow work. If you ask for my involvement in university business that's not absolutely necessary, I might respond with a defense I learned from the department chair who hired me: "Talk to me after tenure." Another tactic that works well for me is to be clear in my refusal but ambiguous in my explanation for the refusal. The key is to avoid providing enough specificity about the excuse that the requester has the opportunity to defuse it. If, for example, I turn down a time-consuming speaking

invitation with the excuse that I have other trips scheduled for around the same time, I don't provide details—which might leave the requester the ability to suggest a way to fit his or her event into my existing obligations—but instead just say, "Sounds interesting, but I can't make it due to schedule conflicts." In turning down obligations, I also resist the urge to offer a consolation prize that ends up devouring almost as much of my schedule (e.g., "Sorry I can't join your committee, but I'm happy to take a look at some of your proposals as they come together and offer my thoughts"). A clean break is best.

In addition to carefully guarding my

obligations, I'm incredibly conscientious about managing my time. Because my time is limited each day, I cannot afford to allow a large deadline to creep up on me, or a morning to be wasted on something trivial, because I didn't take a moment to craft a smart plan. The Damoclean cap on the workday enforced by fixed-schedule productivity has a way of keeping my organization efforts sharp. Without this looming cutoff, I'd likely end up more lax in my habits.

To summarize these observations, Nagpal and I can both succeed in academia without Tom-style overload due to two reasons. First, we're asymmetric in the culling forced by our

fixed-schedule commitment. By ruthlessly reducing the shallow while preserving the deep, this strategy frees up our time without diminishing the amount of new value we generate. Indeed, I would go so far as to argue that the reduction in shallow frees up *more* energy for the deep alternative, allowing us to produce more than if we had defaulted to a more typical crowded schedule. Second, the limits to our time necessitate more careful thinking about our organizational habits, also leading to more value produced as compared to longer but less organized schedules.

The key claim of this strategy is that these same benefits hold for most

knowledge work fields. That is, even if you're not a professor, fixed-schedule productivity can yield powerful benefits. In most knowledge work jobs, it can be difficult in the moment to turn down a shallow commitment that seems harmless in isolation—be it accepting an invitation to get coffee or agreeing to "jump on a call." A commitment to fixed-schedule productivity, however, shifts you into a scarcity mind-set. Suddenly any obligation beyond your deepest efforts is suspect and seen as potentially disruptive. Your default answer becomes no, the bar for gaining access to your time and attention rises precipitously, and you begin to organize

the efforts that pass these obstacles with a ruthless efficiency. It might also lead you to test assumptions about your company's work culture that you thought were ironclad but turn out to be malleable. It's common, for example, to receive e-mails from your boss after hours. Fixed-schedule productivity would have you ignore these messages until the next morning. Many suspect that this would cause problems, as such responses are *expected*, but in many cases, the fact that your boss happens to be clearing her inbox at night doesn't mean that she expects an immediate response—a lesson this strategy would soon help you discover.

Fixed-schedule productivity, in other words, is a *meta-habit* that's simple to adopt but broad in its impact. If you have to choose just one behavior that reorients your focus toward the deep, this one should be high on your list of possibilities. If you're still not sure, however, about the idea that artificial limits on your workday can make you more successful, I urge you to once again turn your attention to the career of fixed-schedule advocate Radhika Nagpal. In a satisfying coincidence, at almost the exact same time that Tom was lamenting online about his unavoidably intense workload as a young professor, Nagpal was celebrating the latest of the many

professional triumphs she has experienced despite her fixed schedule: Her research was featured on the cover of the journal *Science*.

Become Hard to Reach

No discussion of shallow work is complete without considering e-mail. This quintessential shallow activity is particularly insidious in its grip on most knowledge workers' attention, as it delivers a steady stream of distractions *addressed specifically to you.* Ubiquitous e-mail access has become so ingrained in our professional habits that

we're beginning to lose the sense that we have any say in its role in our life. As John Freeman warns in his 2009 book, *The Tyranny of E-mail*, with the rise of this technology "we are slowly eroding our ability to explain—in a careful, complex way—why it is so wrong for us to complain, resist, or redesign our workdays so that they are manageable." E-mail seems a fait accompli. Resistance is futile.

This strategy pushes back at this fatalism. Just because you cannot avoid this tool altogether doesn't mean you have to cede all authority over its role in your mental landscape. In the following sections I describe three tips that will

help you regain authority over how this technology accesses your time and attention, and arrest the erosion of autonomy identified by Freeman. Resistance is not futile: You have more control over your electronic communication than you might at first assume.

***Tip #1**: Make People Who Send You E-mail Do More Work*

Most nonfiction authors are easy to reach. They include an e-mail address on their author websites along with an

open invitation to send them any request or suggestion that comes to mind. Many even encourage this feedback as a necessary commitment to the elusive but much-touted importance of "community building" among their readers. But here's the thing: *I don't buy it*.

If you visit the contact page on my author website, there's no general-purpose e-mail address. Instead, I list different individuals you can contact for specific purposes: my literary agent for rights requests, for example, or my speaking agent for speaking requests. If you want to reach me, I offer only a special-purpose e-mail address that comes with conditions and a lowered

expectation that I'll respond:

> *If you have an offer, opportunity, or introduction that might make my life more interesting, e-mail me at* ***interesting [at] calnewport.com****. For the reasons stated above, I'll only respond to those proposals that are a good match for my schedule and interests.*

I call this approach a *sender filter*, as I'm asking my correspondents to filter themselves before attempting to contact me. This filter has significantly reduced the time I spend in my inbox. Before I

began using a sender filter, I had a standard general-purpose e-mail address listed on my website. Not surprisingly, I used to receive a large volume of long e-mails asking for advice on specific (and often quite complicated) student or career questions. I like to help individuals, but these requests became overwhelming—they didn't take the senders long to craft but they would require a lot of explanation and writing on my part to respond. My sender filter has eliminated most such communication, and in doing so, has drastically reduced the number of messages I encounter in my writing inbox. As for my own interest in helping

my readers, I now redirect this energy toward settings I carefully choose to maximize impact. Instead of allowing any student in the world to send me a question, for example, I now work closely with a small number of student groups where I'm quite accessible and can offer more substantial and effective mentoring.

Another benefit of a sender filter is that it resets expectations. The most crucial line in my description is the following: "I'll only respond to those proposals that are a good match for my schedule and interests." This seems minor, but it makes a substantial difference in how my correspondents

think about their messages to me. The default social convention surrounding e-mail is that unless you're famous, if someone sends you something, you owe him or her a response. For most, therefore, an inbox full of messages generates a major sense of obligation.

By instead resetting your correspondents' expectations to the reality that you'll probably *not* respond, the experience is transformed. The inbox is now a collection of opportunities that you can glance at when you have the free time—seeking out those that make sense for you to engage. But the pile of unread messages no longer generates a sense of obligation. You could, if you wanted to,

ignore them all, and nothing bad would happen. Psychologically, this can be freeing.

I worried when I first began using a sender filter that it would seem pretentious—as if my time was more valuable than that of my readers—and that it would upset people. But this fear wasn't realized. Most people easily accept the idea that you have a right to control *your own* incoming communication, as they would like to enjoy this same right. More important, people appreciate clarity. Most are okay to not receive a response if they don't expect one (in general, those with a minor public presence, such as authors,

overestimate how much people really care about their replies to their messages).

In some cases, this expectation reset might even earn you *more* credit when you do respond. For example, an editor of an online publication once sent me a guest post opportunity with the assumption, set by my filter, that I would likely not respond. When I did, it proved a happy surprise. Here's her summary of the interaction:

> *So, when I emailed Cal to ask if he wanted to contribute to [the publication], my expectations were set. He didn't have*

anything on his [sender filter] about wanting to guest blog, so there wouldn't have been any hard feelings if I'd never heard a peep. Then, when he did respond, I was thrilled.

My particular sender filter is just one example of this general strategy. Consider consultant Clay Herbert, who is an expert in running crowd-funding campaigns for technology start-ups: a specialty that attracts a lot of correspondents hoping to glean some helpful advice. As a Forbes.com article on sender filters reports, "At some point, the number of people reaching out

exceeded [Herbert's] capacity, so he created filters that put the onus on the person asking for help."

Though he started from a similar motivation as me, Herbert's filters ended up taking a different form. To contact him, you must first consult an FAQ to make sure your question has not already been answered (which was the case for a lot of the messages Herbert was processing before his filters were in place). If you make it through this FAQ sieve, he then asks you to fill out a survey that allows him to further screen for connections that seem particularly relevant to his expertise. For those who make it past this step, Herbert enforces a

small fee you must pay before communicating with him. This fee is not about making extra money, but is instead about selecting for individuals who are serious about receiving and acting on advice. Herbert's filters still enable him to help people and encounter interesting opportunities. But at the same time, they have reduced his incoming communication to a level he can easily handle.

To give another example, consider Antonio Centeno, who runs the popular *Real Man Style* blog. Centeno's sender filter lays out a two-step process. If you have a question, he diverts you to a public location to post it. Centeno thinks

it's wasteful to answer the same questions again and again in private one-on-one conversations. If you make it past this step, he then makes you commit to, by clicking check boxes, the following three promises:

- ✓ I am not asking Antonio a style question I could find searching Google for 10 minutes.
- ✓ I am not SPAMMING Antonio with a cut-and-pasting generic request to promote my unrelated business.
- ✓ I will do a good deed for some random stranger if Antonio

responds within 23 hours.

The message box in which you can type your message doesn't appear on the contact page until after you've clicked the box by all three promises.

To summarize, the technologies underlying e-mail are transformative, but the current social conventions guiding how we apply this technology are underdeveloped. The notion that all messages, regardless of purpose or sender, arrive in the same undifferentiated inbox, and that there's an expectation that every message deserves a (timely) response, is absurdly unproductive. The sender filter

is a small but useful step toward a better state of affairs, and is an idea whose time has come—at least for the increasing number of entrepreneurs and freelancers who both receive a lot of incoming communication and have the ability to dictate their accessibility. (I'd also love to see similar rules become ubiquitous for intra-office communication in large organizations, but for the reasons argued in Chapter 2, we're probably a long way from that reality.) If you're in a position to do so, consider sender filters as a way of reclaiming some control over your time and attention.

***Tip #2**: Do More Work When* You *Send or Reply to E-mails*

Consider the following standard e-mails:

E-mail #1: "It was great to meet you last week. I'd love to follow up on some of those issues we discussed. Do you want to grab coffee?"

E-mail #2: "We should get back to the research problem we discussed during my last visit. Remind me where we are with that?"

E-mail #3: "I took a stab at that article we discussed. It's attached. Thoughts?"

These three examples should be familiar to most knowledge workers, as they're representative of many of the messages that fill their inboxes. They're also potential productivity land mines: How you respond to them will have a significant impact on how much time and attention the resulting conversation ultimately consumes.

In particular, interrogative e-mails like these generate an initial instinct to dash off the quickest possible response that will clear the message—temporarily—out of your inbox. A quick response

will, in the short term, provide you with some minor relief because you're bouncing the responsibility implied by the message off your court and back onto the sender's. This relief, however, is short-lived, as this responsibility will continue to bounce back again and again, continually sapping your time and attention. I suggest, therefore, that the right strategy when faced with a question of this type is to pause a moment before replying and take the time to answer the following key prompt:

> *What is the* project *represented by this message, and what is the most efficient (in terms of*

messages generated) process *for bringing this project to a successful conclusion?*

Once you've answered this question for yourself, replace a quick response with one that takes the time to describe the process you identified, points out the current step, and emphasizes the step that comes next. I call this the *process-centric approach* to e-mail, and it's designed to minimize both the number of e-mails you receive and the amount of mental clutter they generate.

To better explain this process and why it works consider the following process-centric responses to the sample

e-mails from earlier:

Process-Centric Response to E-mail #1: "I'd love to grab coffee. Let's meet at the Starbucks on campus. Below I listed two days next week when I'm free. For each day, I listed three times. If any of those day and time combinations work for you, let me know. I'll consider your reply confirmation for the meeting. If none of those date and time combinations work, give me a call at the number below and we'll hash out a time that works. Looking forward to it."

Process-Centric Response to E-mail #2: "I agree that we should return to this problem. Here's what I suggest…

"Sometime in the next week e-mail me everything you remember about our discussion on the problem. Once I receive that message, I'll start a shared directory for the project and add to it a document that summarizes what you sent me, combined with my own memory of our past discussion. In the document, I'll highlight the two or three most promising next steps.

"We can then take a crack at those next steps for a few weeks and check back in. I suggest we schedule a phone call for a month from now for this purpose. Below I listed some dates and times when I'm available for a call. When you respond with your notes,

indicate the date and time combination that works best for you and we'll consider that reply confirmation for the call. I look forward to digging into this problem."

Process-Centric Response to E-mail #3: "Thanks for getting back to me. I'm going to read this draft of the article and send you back an edited version annotated with comments on Friday (the 10th). In this version I send back, I'll edit what I can do myself, and add comments to draw your attention to places where I think you're better suited to make the improvement. At that point, you should have what you need to polish and submit the final draft, so I'll leave

you to do that—no need to reply to this message or to follow up with me after I return the edits—unless, of course, there's an issue."

In crafting these sample responses, I started by identifying the project implied by the message. Notice, the word "project" is used loosely here. It can cover things that are large and obviously projects, such as making progress on a research problem (Example #2), but it applies just as easily to small logistical challenges like setting up a coffee meeting (Example #1). I then took a minute or two to think through a process that gets us from the current state to a

desired outcome with a minimum of messages required. The final step was to write a reply that clearly describes this process and where we stand. These examples centered on an e-mail reply, but it should be clear that a similar approach also works when writing an e-mail message from scratch.

The process-centric approach to e-mail can significantly mitigate the impact of this technology on your time and attention. There are two reasons for this effect. First, it reduces the number of e-mails in your inbox—sometimes significantly (something as simple as scheduling a coffee meeting can easily spiral into half a dozen or more

messages over a period of many days, if you're not careful about your replies). This, in turn, reduces the time you spend in your inbox and reduces the brainpower you must expend when you do.

Second, to steal terminology from David Allen, a good process-centric message immediately "closes the loop" with respect to the project at hand. When a project is initiated by an e-mail that you send or receive, it squats in your mental landscape—becoming something that's "on your plate" in the sense that it has been brought to your attention and eventually needs to be addressed. This method closes this open loop as soon as

it forms. By working through the whole process, adding to your task lists and calendar any relevant commitments on your part, and bringing the other party up to speed, your mind can reclaim the mental real estate the project once demanded. Less mental clutter means more mental resources available for deep thinking.

Process-centric e-mails might not seem natural at first. For one thing, they require that you spend more time thinking about your messages before you compose them. In the moment, this might seem like you're spending *more* time on e-mail. But the important point to remember is that the extra two to three

minutes you spend at this point will save you many more minutes reading and responding to unnecessary extra messages later.

The other issue is that process-centric messages can seem stilted and overly technical. The current social conventions surrounding e-mail promote a conversational tone that clashes with the more systematic schedules or decision trees commonly used in process-centric communication. If this concerns you, I suggest that you add a longer conversational preamble to your messages. You can even separate the process-centric portion of the message from the conversational opening with a

divider line, or label it "Proposed Next Steps," so that its technical tone seems more appropriate in context.

In the end, these minor hassles are worth it. By putting more thought up front into what's really being proposed by the e-mail messages that flit in and out of your inbox, you'll greatly reduce the negative impact of this technology on your ability to do work that actually matters.

Tip #3: Don't Respond

As a graduate student at MIT, I had the opportunity to interact with famous

academics. In doing so, I noticed that many shared a fascinating and somewhat rare approach to e-mail: Their default behavior when receiving an e-mail message is to *not* respond.

Over time, I learned the philosophy driving this behavior: When it comes to e-mail, they believed, it's the sender's responsibility to convince the receiver that a reply is worthwhile. If you didn't make a convincing case *and* sufficiently minimize the effort required by the professor to respond, you didn't get a response.

For example, the following e-mail would likely not generate a reply with many of the famous names at the

Institute:

> *Hi professor. I'd love to stop by sometime to talk about <topic X>. Are you available?*

Responding to this message requires too much work ("Are you available?" is too vague to be answered quickly). Also, there's no attempt to argue that this chat is worth the professor's time. With these critiques in mind, here's a version of the same message that would be more likely to generate a reply:

> *Hi professor. I'm working on a project similar to <topic X>*

with my advisor, <professor Y>. Is it okay if I stop by in the last fifteen minutes of your office hours on Thursday to explain what we're up to in more detail and see if it might complement your current project?

Unlike the first message, this one makes a clear case for why this meeting makes sense and minimizes the effort needed from the receiver to respond.

This tip asks that you replicate, to the extent feasible in your professional context, this professorial ambivalence to e-mail. To help you in this effort, try applying the following three rules to sort

through which messages require a response and which do not.

Professorial E-mail Sorting: Do not reply to an e-mail message if any of the following applies:

- It's ambiguous or otherwise makes it hard for you to generate a reasonable response.
- It's not a question or proposal that interests you.
- Nothing really good would happen if you did respond and nothing really bad would happen if you didn't.

In all cases, there are many obvious exceptions. If an ambiguous message about a project you don't care about comes from your company's CEO, for example, you'll respond. But looking beyond these exceptions, this professorial approach asks you to become way more ruthless when deciding whether or not to click "reply."

This tip can be uncomfortable at first because it will cause you to break a key convention currently surrounding e-mail: Replies are assumed, regardless of the relevance or appropriateness of the message. There's also no way to avoid that some bad things will happen if you take this approach. At the minimum,

some people might get confused or upset—especially if they've never seen standard e-mail conventions questioned or ignored. Here's the thing: This is okay. As the author Tim Ferriss once wrote: "Develop the habit of letting small bad things happen. If you don't, you'll never find time for the life-changing big things." It should comfort you to realize that, as the professors at MIT discovered, people are quick to adjust their expectations to the specifics of your communication habits. The fact you didn't respond to their hastily scribed messages is probably not a central event in their lives.

Once you get past the discomfort of

this approach, you'll begin to experience its rewards. There are two common tropes bandied around when people discuss solutions to e-mail overload. One says that sending e-mails generates more e-mails, while the other says that wrestling with ambiguous or irrelevant e-mails is a major source of inbox-related stress. The approach suggested here responds aggressively to both issues—you send fewer e-mails and ignore those that aren't easy to process—and by doing so will significantly weaken the grip your inbox maintains over your time and attention.